



Music Emotion Recognition (MER)

Sihan, Yanxin, Qianyu

Agenda

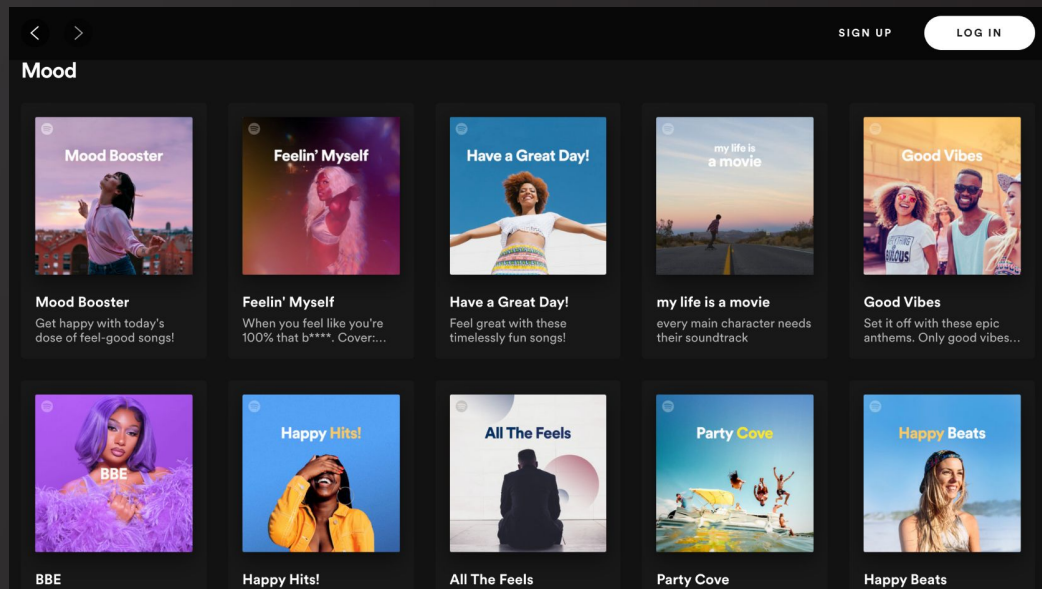
1. Project Introduction
2. Data Crawling & Preprocessing
3. Feature vectors & General Model(SVM, Tree, etc.) Comparison
4. Audio feature and Convolutional Neural Network Analysis

MER

- multi-label tagging task
- acoustic features \leftrightarrow emotion tags
- application: music recommendation etc.
- models: SVM, random forests, CNN
(Liu, Xin, et al.)

Dataset -- overview

- source: Spotify
- 7195 songs (5510 with .wav)
 - Calm: 2951 (2742 with .wav)
 - sad: 439 (232 with .wav)
 - energetic: 3554 (2364 with .wav)
 - happy: 251 (172 with .wav)
- 16 features
 - provided by Spotify API
- Label
 - self-labelled
 - binary: positive & negative
 - multi-class



Dataset -- preprocessing

- drop duplicated songs
- extract songs with preview_url (with audio)


```
calm.csv  
calm_no_duplicates.csv  
calm_no_duplicates_no_nan.csv  
calm_vectors/
```

Dataset -- features


- object metadata
 - 16 features from Spotify API

```
id                                0LweQRsfJ3pRAJJFy6DrR1
danceability                      0.776
energy                           0.692
key                               2
loudness                         -3.834
mode                             1
speechiness                      0.0555
acousticness                    0.0075
instrumentalness                 0
liveness                        0.0588
valence                         0.543
tempo                           95.972
duration_ms                      198174
time_signature                   4
artist_name                      Jessi
track_name                       What Type of X
preview_url                      https://p.scdn.co/mp3-preview/917623a4638d1005...
binary_label                     1
multiclass_label                 4
Name: 1, dtype: object
```


Dataset-features



calm



energetic

Spectrogram

- **2D discrete Fourier transform of audio**
- **visualize harmonic structure in audio**

Dataset -- features



- 30 second preview
 - retrieve .wav
 - .wav -> matrix $128 * 216$
 - intensities of spaced frequency ranges over discrete slices in time

```
x, sr = librosa.load(file, offset = 10, duration = 5)
ps = librosa.feature.melspectrogram(x, sr = sr)
```

```
data = np.array(data)
print(data.shape)
np.save(name + '.npy', data)
```



Traditional models :)

SVM, Random Forest, etc.



Final Result

	precision	recall	f1-score	support
Sad	0.43	0.48	0.45	88
Healing	0.95	0.97	0.96	1926
joyful	0.90	0.86	0.88	816

first vs. last

	precision	recall	f1-score	support
Sad	0.43	0.48	0.45	88
Healing	0.95	0.97	0.96	1926
joyful	0.90	0.86	0.88	816

classifier = stacking (SVM, Decision Tree, Random Forest)

	precision	recall	f1-score	support
sad	1.00	0.06	0.11	53
calm	0.93	1.00	0.96	1926
happy	1.00	0.02	0.04	50
energetic	0.89	0.85	0.87	766

classifier = SVM



Model Improvement

-----A long journey of model optimization-----



Better data

- oversampling
- add more data

Better model

- grid search for hyperparameter
- try different models
- stack models

Other

- change classification criterion

Improve the model

Main Difficulty: class imbalance

calm	:	9629	68.91%	
energetic	:	3829		27.40%
sad	:	266	1.90%	
happy	:	250	1.79%	

Better Data

Oversampling: library SMOTE

before	precision	recall	f1-score	support
sad	1.00	0.06	0.11	53
calm	0.93	1.00	0.96	1926
happy	1.00	0.02	0.04	50
energetic	0.89	0.85	0.87	766

classifier = SVM default parameters

after	precision	recall	f1-score	support
sad	0.17	0.62	0.27	53
calm	0.94	0.95	0.94	1926
happy	0.15	0.58	0.23	50
energetic	0.94	0.58	0.72	766

classifier = SVM default parameters + oversampling

More data helps ! ! !

	precision	recall	f1-score	support
sad	0.33	0.72	0.45	88
calm	0.94	0.96	0.95	1926
happy	0.16	0.58	0.25	50
energetic	0.94	0.62	0.75	766

SVM default parameters + oversampling + add data

Better Model

Hyperparameter tuning: Gridsearch

```
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import make_scorer, f1_score
param_grid = {'C': [1, 10, 100, 1000],
              'gamma': [1, 0.1, 0.01, 0.001],
              'kernel': ['rbf']}

svc_clf = SVC()
scorer = make_scorer(f1_score, average = 'weighted')
grid = GridSearchCV(SVC(), param_grid, scoring=scorer, refit = True)
```


Try different models:

```
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
```

Stacking

```
models = [('dtree', dtree),
          ('rforest', rforest),
          ('svm', svm)]
stacking = StackingClassifier(
    estimators=models)
```

Other



Change classification criterion

Future Work

- **More data : add data for the minority class**
- **Data of higher quality**
- **More features: extract features from the audio**

Neural Network

Feature vector input:

- One-layer NN


Raw audio input:

- Conv + FC
- Conv + LSTM + DNN


Simple NN

One layer with 8 nodes

Accuracy Score: 89.4%



```
In [351]: #Normalize the features  
X = MinMaxScaler().fit_transform(X)
```



Model: "sequential_46"


Layer (type)	Output Shape	Param #
=====		
dense_92 (Dense)	(None, 8)	112
=====		
dense_93 (Dense)	(None, 4)	36
=====		
Total params: 148		
Trainable params: 148		
Non-trainable params: 0		

Baseline model:


Three-layer CNN

Accuracy: 53.9%

Conv + FC



Conv + LSTM + DNN



(128, 216, 1)

Conv + FC

3,552,513

Parameters


30+ mins

Training time

81%


Best Val
Accuracy

Can we do better??



Adam Optimizer

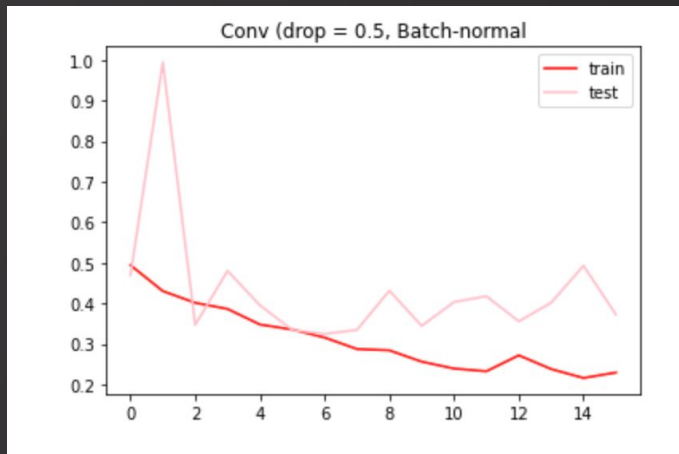
Cross-entropy Loss



Model Optimization

Callbacks=[es]
Early Stopping

Batch size
32



Batch-normalization

After Activation
Before drop out


Drop Out rate
0.3 --> 0.5

Accuracy
89.47%


```
es = EarlyStopping(monitor='val_loss', mode='min', verbose=1, patience=3)
```

Can we do better?

We changed the filter size
from square
into rectangular...



```
model = Sequential()
```



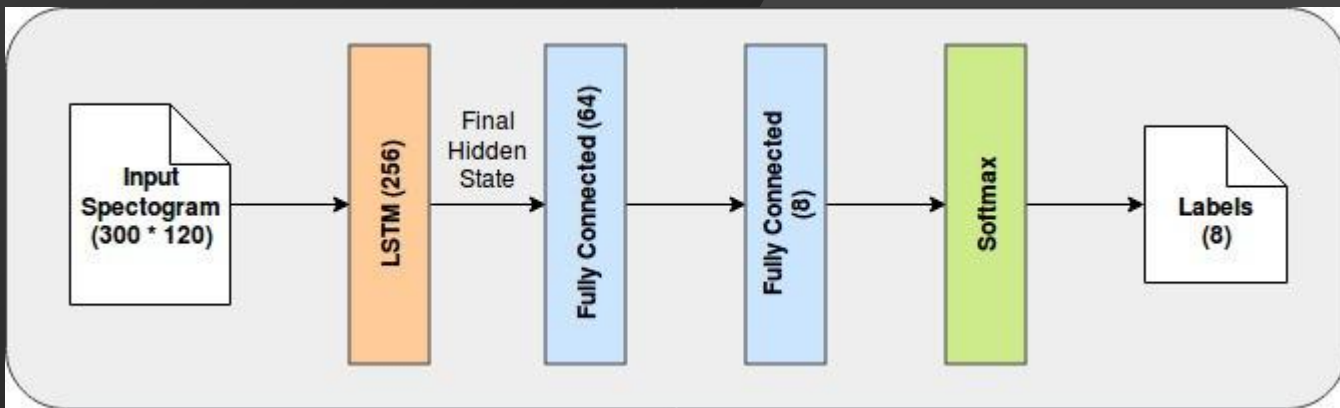
```
model.add(Dense(100, activation='relu'))
```

With only 446193 parameters!
1/10 of original numbers!

Accuracy: 89.11%

Pure LSTM

Low training rate
and bad behavior!



Conv + LSTM + DNN

Filter

kernel_size=(1, 3)

Time + LSTM

Bidirectional LSTM

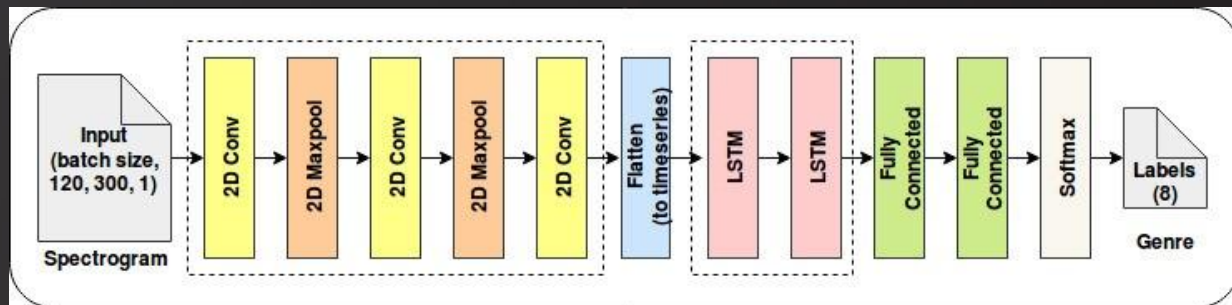
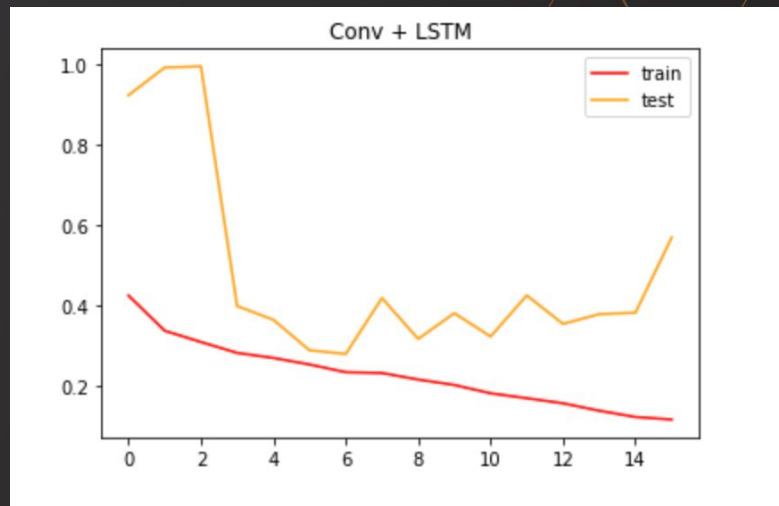
(128, dropout=0.25,
return_sequences=True)

89.11%

High Training Accuracy !

LeakyReLU(alpha=0.01)

10+ mins/Epoch ...
Super slow ... Orz ...



Model Limitation & Future Improvement

- 1. Data limitation**
 - a. Subjective label / 5s Clips / quantity
- 2. Audio Data Augmentation**
 - a. Time Stretching / Pitch Shifting
 - b. Dynamic Range Compression
 - c. Background Noise
- 3. Multi-class Emotion Recognition**
- 4. New data type (SoundNet)**



Thank you!

